

Xian Wu 903142009

Question 4:

My Computer:

```
lisa@lisaMbp:~/Downloads/hw1$ lscpu
Architecture:          x86_64
CPU op-mode(s):        32-bit, 64-bit
Byte Order:            Little Endian
CPU(s):                4
On-line CPU(s) list:  0-3
Thread(s) per core:   2
Core(s) per socket:   2
Socket(s):             1
NUMA node(s):          1
Vendor ID:             GenuineIntel
CPU family:            6
Model:                 61
Stepping:              4
CPU MHz:               1410.804
BogoMIPS:              5799.82
Virtualization:        VT-x
L1d cache:             32K
L1i cache:             32K
L2 cache:               256K
L3 cache:               3072K
NUMA node0 CPU(s):    0-3
```

And this is all my code revised:

```
if(threads == 1) {
    time_spent=0;
    for(int count=0;count<20;count++){
        start = omp_get_wtime();
        sorted = bubble_sort_serial(unsorted,num_elems);
        end = omp_get_wtime();
        timer[count]=(double)(end-start);
        time_spent+=(double)(end-start);
    }
    time_spent/=20.0;
} else {
    time_spent=0;
    for(int count=0;count<20;count++){
        start = omp_get_wtime();
        sorted = bubble_sort_parallel(unsorted,num_elems,threads);
        end = omp_get_wtime();
        timer[count]=(double)(end-start);
        time_spent+=(double)(end-start);
    }
}
```

```

    time_spent/=20.0;
}

double tmp=0.0;
for(int i=0;i<20;i++){
    tmp+=pow(timer[i]-time_spent,2);
}
tmp=pow(tmp,0.5)/time_spent;
printf("variance %lg \n",tmp);
printf("Time Spent %lg \n",time_spent);
*****

```

Fluctuation:

I calculated the average of 20 times experiment to get each value in the table to reduce the fluctuation effect. Also I calculated the variance for each problem size. If (variance/average) is too large, the set of data for this problem size could be considered unreliable.

P \ N	1000	10000	20000	40000	60000	80000
1	0.00221152	0.238954	1.04162	4.38342	10.2365	17.9679
2	0.00214644	0.149752	0.54677	2.77651	5.17902	9.88461
4	0.00252255	0.12589	0.506221	2.46537	4.99839	9.72908
6	0.0100587	0.223523	0.793997			
8	0.0136448	0.275474	0.908305			

For problem size 1000, the variance/average ratio could be from 1.5 to 2.5, so we should not consider the 1000-column data reliable. There is serious fluctuation.

Strong Scaling

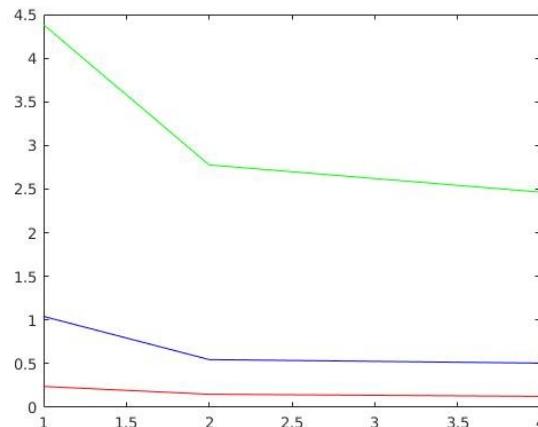
By keeping the problem size and varying the number of threads:

You can see clearly that within **number of threads no more than 4**, the more threads used the less runtime costed. But runtime does not reduces linearly, which is due to the overhead for processors to communicate and synchronize.

Also as for why a thread number larger than 4 fails to make it faster, it is because my computer has just 4 CPU cores, if you uses more than 4 threads, CPUs have to do context switch to simulate the multi-processor situation, which exactly increases the latency.

The speedup would be

See the graph below:



The x-label is number of processors and the y-label is the time spent.

Weak Scaling

By keeping the processor size proportional to the problem size and calculating the speedup:

P \ N	10000	20000	40000
1	0.238954	1.04162	4.38342
2		0.54677	
4			2.46537
speedup	--	1.9050	1.7780

Since there is not enough data for plotting a graph, but you can still see a speedup around 1.9.

Here the sequential part of the program is removed out of consideration. So in the ideal case the speedup should be :

$$\text{speedup} = N \text{ times the problem size grows}$$

However, here about the bubble sort algorithm, with single processor, when the problem size grows, the runtime does not increase linearly, which made the speedup of 40000-size large problem even less than 2.